“A Hybrid Deep Learning and Machine Learning Framework for Fruit Damage Segmentation and Quality Assessment”

Department of Information Systems and Business Analytics

Kent State University

Chaojiang (CJ) Wu, Ph.D.

Submitted by:

Shivani Bandari

bshivani@kent.edu

811363266

INTRODUCTION

- Fruit quality inspection is important for preventing spoiled or damaged fruits from reaching consumers.
- Manual inspection is slow, inconsistent, and prone to human error.
- Deep learning allows computers to automatically analyze fruit images and detect damage accurately.
- This project uses computer vision to classify fruits as healthy or damaged.
- A combination of CNN (Deep Learning) and Machine Learning models is used for comparison.
- The goal is to build a reliable and efficient system for automated fruit quality assessment.

MOTIVATION

- Damaged fruits often go unnoticed during manual inspection, leading to customer complaints and economic loss.
- Human inspection is slow, inconsistent, and depends on individual experience.
- Automated detection using deep learning provides faster, more accurate, and repeatable results.
- Computer vision can identify subtle damage patterns that may not be visible to the human eye.
- Using AI for fruit quality assessment helps improve food safety, reduce waste, and support large-scale agricultural operations.
- This project aims to show how ML and DL can solve real-world problems in agriculture.

DATASET DESCRIPTION

- The dataset contains high-quality images of fruits labeled as healthy or damaged.
- Images show variations in lighting, color, texture, and different types of damage.
- The dataset is suitable for both classification and segmentation tasks.
- All images were resized, normalized, and augmented before training.
- The data was divided into training, validation, and testing sets for fair evaluation.

METHODOLOGY

- Collect fruit images containing both healthy and damaged samples.
- Preprocess the images by resizing, normalizing, and applying data augmentation.
- Train a U-Net model to segment and highlight the damaged regions on the fruit.
- Develop a CNN classifier to categorize images as healthy or damaged.
- Build machine learning baseline models (SVM, Logistic Regression, Random Forest) for comparison.
- Evaluate all models using accuracy, loss curves, confusion matrices, and visual outputs.
- Interpret the results through segmentation masks, predictions, and performance graphs.

LOSS CURVE

- This graph shows how the training loss and validation loss change during each epoch.
- Loss measures how well the model is learning and improving over time.
- A decreasing loss means the model is learning the patterns in the images correctly.
- Both training and validation loss reduced steadily, indicating stable and effective learning.
- The similar trend between the two curves shows that the model did not overfit.

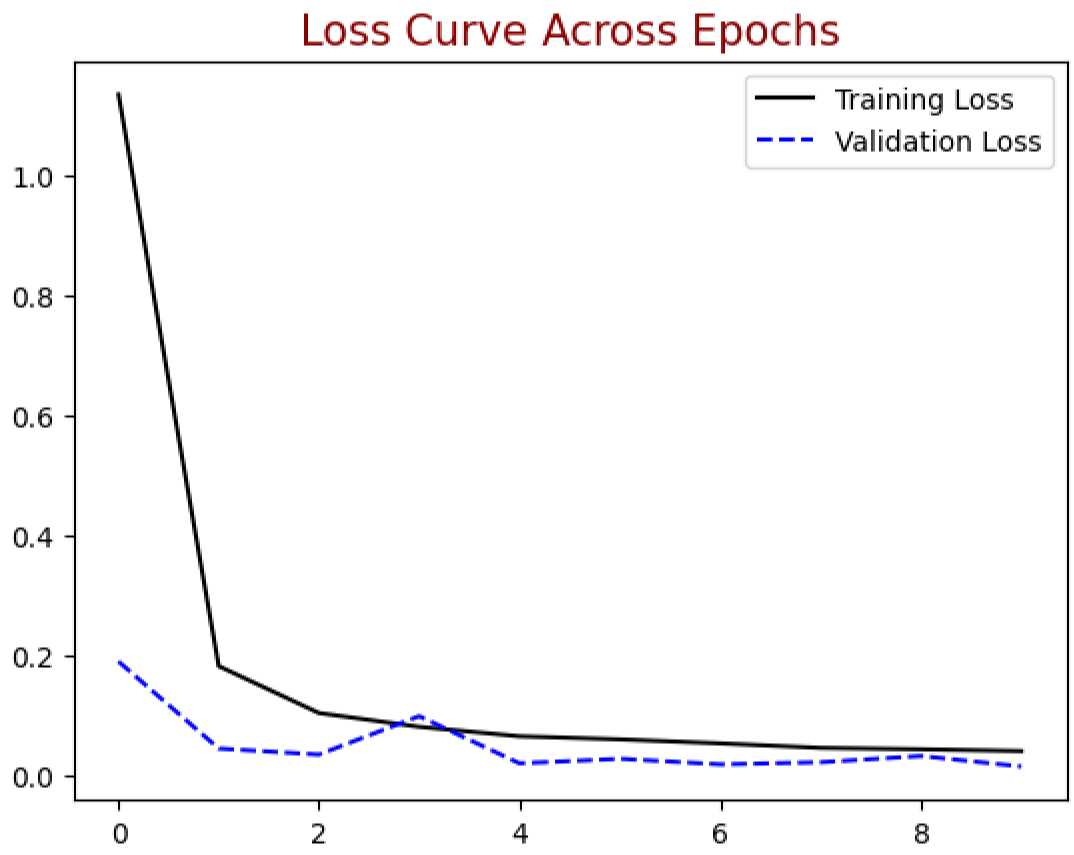

ACCURACY CURVE

- This graph shows how the training accuracy and validation accuracy improved across epochs.
- Accuracy indicates how correctly the model classified healthy and damaged fruit images.
- Both curves steadily increased and then stabilized, showing effective learning.
- The close alignment between training and validation accuracy suggests the model generalizes well.
- The final accuracy levels indicate that the CNN model can reliably classify new fruit images.

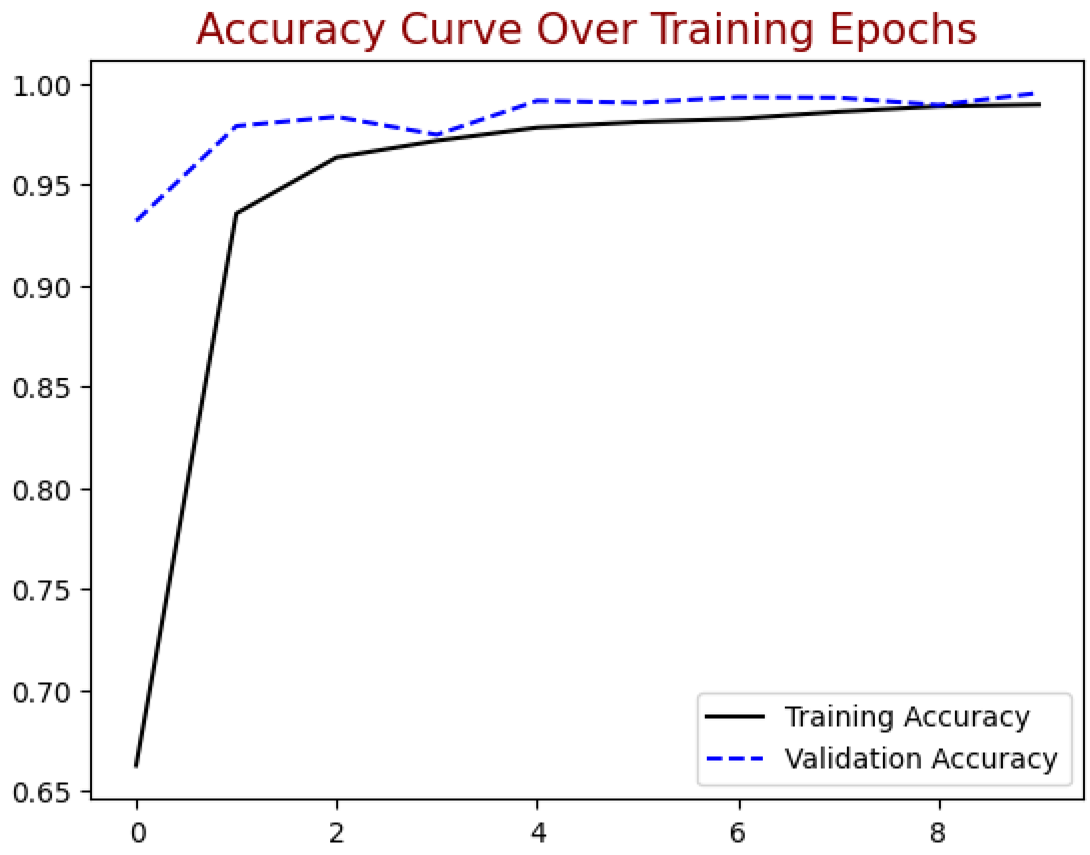

RESULTS

- The CNN model successfully classified fruit images as healthy or damaged.
- Training and validation accuracy increased and stabilized, showing strong learning.
- Training and validation loss decreased, indicating efficient optimization.
- **Final Training Accuracy: 98.99%**
- **Final Validation Accuracy: 99.53%**
- **Final Training Loss: 0.0385**
- **Final Validation Loss: 0.0138**
- Model predictions on test images matched the actual fruit condition in most cases.

Overall, the deep learning approach proved highly effective for automated fruit damage detection.

CONCLUSION

- The CNN model accurately distinguished healthy and damaged fruits, achieving strong performance on both training and validation sets.
- The model reached a training accuracy of 98.99% and validation accuracy of 99.53%, showing excellent generalization.
- Loss values decreased steadily, confirming stable learning and proper optimization.
- Predictions on test images closely matched real fruit conditions, proving the model's reliability.
- Overall, the deep learning framework provides an effective and automated solution for fruit damage detection and quality assessment.

FUTURE WORK

- Expanding the dataset with more fruit types and a wider range of damage patterns.
- Implementing advanced deep learning models such as Vision Transformers (ViT), EfficientNet, or YOLO detectors for higher accuracy.
- Integrating U-Net or other segmentation models for even more precise damage localization.
- Deploying the model on mobile or edge devices for real-time fruit quality inspection.
- Exploring additional data sources like hyperspectral or 3D imaging to detect internal or hidden damage.
- Applying hyperparameter tuning and ensemble methods to further improve robustness and performance.